

An Introduction to RNA-sequencing

Differential Expression Analysis

W. Evan Johnson, Ph.D.
Professor, Division of Infectious Disease
Director, Center for Data Science
Co-Director, Center for Biomedical Informatics and Health AI
Rutgers University – New Jersey Medical School

2026-06-16

Installing R Packages:

Install the following tools for differential expression analysis:
SummarizedExperiment, edgeR, DESeq2, limma,
ComplexHeatmap, TBSignatureProfiler, and umap. We will also
need help from the tidyverse.

```
if (!requireNamespace("BiocManager", quietly = TRUE))  
  install.packages("BiocManager")  
BiocManager::install(c("tidyverse", "SummarizedExperiment",  
                        "edgeR", "DESeq2", "limma", "ComplexHeatmap",  
                        "TBSignatureProfiler", "umap"))
```

Load Packages for Today

We will be using the following packages for our differential expression lecture:

```
library(tidyverse) ## tools for data wrangling
library(SummarizedExperiment) ## managing counts data
library(edgeR) ## differential expression
library(DESeq2) ## differential expression
library(limma) ## differential expression
library(ComplexHeatmap) ## Heatmap visualization
library(TBSignatureProfiler) ## TB signature analysis
library(umap) ## dimension reduction and plotting data
```

Picking Up From Alignment

In the previous lecture we built BAMs and a feature count matrix using **Rsubread / STAR + featureCounts**. Today's starting point is a tidy gene-by-sample **counts table** (`features_combined.txt`) plus per-sample **metadata** (`meta_data.txt`).

Everything that follows – normalization, modeling, visualization – assumes that pipeline has already been run.

Data Structures

A data structure is a particular way of organizing data in a computer so that it can be used effectively. The idea is to reduce the space and time complexities of different tasks.

Data Structures

Data structures in R programming are tools for holding multiple values, variables, and sometimes functions

Please think very carefully about the way you manage and store your data! This can make your life much easier and make your code and data cleaner and more portable!

Data Frames

A large proportion of data analysis challenges start with data stored in a data frame. For example, we stored the data for our motivating example in a data frame. You can access this dataset by loading `TBNanostring.rds` object in R:

```
TBNanostring <- readRDS("example_data/TBNanostring.rds")
```

Data Frames

In RStudio we can view the data with the `View` function:

```
View(TBnanosttring)
```


Data Frames


Or in RMarkdown you can use the `datatable` function from the DT package:

```
datatable(TBnanosttring)
```

You will notice that the TB status is found in the first column of the data frame, followed by the genes in the subsequent columns. The rows represent each individual patient.

Advanced Data Structures

There are advanced R data structures that can facilitate object orientated programming. One useful example is the **SummarizedExperiment** object.



Make a SummarizedExperiment

Using an example dataset from: Verma, et al., 2018

```
## read in data
counts <- read.table(
  "example_data/features_combined.txt",
  sep="\t", header=T, row.names=1)
meta_data <- read.table(
  "example_data/hivtb_SRR_metadata.txt",
  sep="\t", header=T, row.names=2)
group <- meta_data$disease_status
meta_data <- meta_data[match(colnames(counts), rownames(meta_data))]
```

Make a SummarizedExperiment

```
## Make SummarizedExperiment  
se_hivtb <- SummarizedExperiment(assays=list(counts=counts),  
                                colData = meta_data)  
  
## Make log counts, counts per million (cpm), logcpm  
se_hivtb <- mkAssay(se_hivtb, log = TRUE,  
                   counts_to_CPM = TRUE)  
  
assays(se_hivtb)
```

```
## List of length 4  
## names(4): counts log_counts counts_cpm log_counts_cpm
```

EdgeR Example

Implements statistical methods for DE analysis based on the negative binomial model:

```
#Gene Filtering
counts<-counts[which(rowSums(counts)>100),]
#Computes library size
dge <- DGEList(counts=counts, group=group)
#TMM normalization
dge <- calcNormFactors(dge)
# Design matrix
design<-model.matrix(~disease_status, data=meta_data)
#Estimates common, trended and tagwise dispersion
dge<-estimateDisp(counts,design)
```

EdgeR Example

In negative binomial models, each gene is given a dispersion parameter. Dispersions control the variances of the gene counts and underestimation will lead to false discovery and overestimation may lead to a lower rate of true discovery.

EdgeR Example

```
# Neg Bin GLM with the dispersion estimates  
fit<-glmFit(counts,design,  
             dispersion=dge$tagwise.dispersion)  
# Performs likelihood ratio test  
# Compares full versus reduced model  
lrt<-glmLRT(fit, coef=2)
```

EdgeR Example

```
# Prints the top results
topTags(lrt)
```

```
## Coefficient:  disease_statusTB-HIV
##              logFC  logCPM      LR      PValue      FDR
## IL1R2      4.334138 8.208316 93.61044 3.841569e-22 6.150736e-18
## AP3B2      5.751207 2.952926 64.01082 1.237377e-15 6.755923e-12
## FCGR1C     2.818442 4.536519 63.96598 1.265865e-15 6.755923e-12
## VNN1       3.150042 8.072140 60.78606 6.362689e-15 2.546825e-11
## CYP1B1     3.135490 6.873383 59.54714 1.193995e-14 3.823412e-11
## IL18R1     2.726093 6.487874 58.88977 1.667570e-14 4.449912e-11
## SOCS3      2.665152 6.476286 55.84118 7.856849e-14 1.797086e-10
## SLC29A1   -4.135726 3.970313 54.67430 1.422541e-13 2.699678e-10
## CACNG8    -4.134739 3.190544 54.35946 1.669722e-13 2.699678e-10
## ZAK       2.601721 6.127149 54.34023 1.686139e-13 2.699678e-10
```


EdgeR Example

```
## Use the a quasi-likelihood F-test  
fit<-glmQLFit(counts, design,  
              dispersion=dge$tagwise.dispersion)  
## use for small datasets  
qlf<-glmQLFTest(fit, coef=2)
```

EdgeR Example

```
# Prints the top results
topTags(qlf)
```

```
## Coefficient:  disease_statusTB-HIV
##              logFC  logCPM      F      PValue      FDR
## IL1R2         4.334144 8.208316 96.16041 1.411235e-22 2.259528e-18
## FCGR1C        2.818462 4.536519 65.27824 7.434147e-16 4.644916e-12
## AP3B2         5.751790 2.952926 64.96483 8.703233e-16 4.644916e-12
## VNN1          3.150045 8.072140 62.45782 3.077241e-15 1.231742e-11
## CYP1B1        3.135482 6.873383 61.22274 5.734734e-15 1.836376e-11
## IL18R1        2.726091 6.487874 60.59190 7.882229e-15 2.103373e-11
## SOCS3         2.665156 6.476286 57.45227 3.842609e-14 8.789144e-11
## ZAK           2.601720 6.127149 55.89059 8.455732e-14 1.621038e-10
## FKBP5         2.426321 7.528467 55.66226 9.489672e-14 1.621038e-10
## SLC29A1       -4.135688 3.970313 55.53406 1.012453e-13 1.621038e-10
```

EdgeR Example

```
#For visualization, heatmaps/PCA  
Logcpm<-cpm(counts,log=TRUE)
```

DESeq2 Example

```
#colData is a data frame of demographic/phenotypic data  
dds <- DESeqDataSetFromMatrix(countData = counts,  
                               colData=meta_data,  
                               design=~disease_status)
```

```
#Gene Filtering  
dds<-dds[rowSums(counts(dds))>100,]
```

DESeq2 Example

```
#Performs estimation of size factors,  
#dispersion, and negative binomial GLM fitting  
dds<-DESeq(dds)
```

```
## estimating size factors
```

```
## Note: levels of factors in the design contain characters other than  
## letters, numbers, '_' and '.'. It is recommended (but not required) to use  
## only letters, numbers, and delimiters '_' or '.', as these are safe characters  
## for column names in R. [This is a message, not a warning or an error]
```

```
## estimating dispersions
```

```
## gene-wise dispersion estimates
```

```
## mean-dispersion relationship
```

```
## Note: levels of factors in the design contain characters other than  
## letters, numbers, '_' and '.'. It is recommended (but not required) to use  
## only letters, numbers, and delimiters '_' or '.', as these are safe characters
```

DESeq2 Example

```
res <- results(dds)[order(results(dds)[,6]),]
res[1:25,]
```

log2 fold change (MLE): disease status TB.HIV.ART vs HIV

Wald test p-value: disease status TB.HIV.ART vs HIV


DataFrame with 25 rows and 6 columns

##	baseMean	log2FoldChange	lfcSE	stat	pvalue	padj
##	<numeric>	<numeric>	<numeric>	<numeric>	<numeric>	<numeric>
## ZEB2	1702.8164	1.75543	0.217589	8.06764	7.16683e-16	1.14748e-11
## DCUN1D3	48.7029	2.52499	0.362090	6.97337	3.09430e-12	2.47714e-08
## TIPARP	724.9429	2.22978	0.333852	6.67894	2.40672e-11	1.28447e-07
## ITPKC	196.8723	3.04144	0.465478	6.53401	6.40298e-11	2.56295e-07
## LAIR1	1528.6852	3.01392	0.471809	6.38800	1.68065e-10	4.25413e-07
## ...	...	...	...	...	...	...
## ATP6V1C1	1006.405	2.50826	0.420512	5.96477	2.44975e-09	1.84789e-06
## TAB3	546.645	1.52153	0.254844	5.97044	2.36615e-09	1.84789e-06
## IL18R1	1777.342	4.19796	0.706351	5.94317	2.79570e-09	1.94617e-06
## NR3C1	2177.200	1.46964	0.247644	5.93451	2.94720e-09	1.96615e-06
## IRAK3	4822.486	4.00584	0.678075	5.90765	3.47012e-09	2.13939e-06

Heatmap of DEGs

```
# Make a Heatmap of DEGs
mat = as.matrix(assay(se_hivtb,"log_counts_cpm")
               ) [order(results(dds)[,6])[1:100],]
               # Using first 100 genes to simplify
mat = t(scale(t(mat))) ## scale gene expression
df=data.frame(Disease=colData(se_hivtb)$disease_status)
ha = HeatmapAnnotation(df = df,
                      col = list(Disease=c(
                        "TB-HIV"="Red",
                        "HIV"="Blue",
                        "TB-HIV-ART"="Green")))
Heatmap(mat,show_row_names=F, show_column_names = F,
        top_annotation = ha)
```

Heatmap of DEGs



Limma Example

- ▶ Most similar to microarray data flow
- ▶ Reads counts are converted to log2 counts per million (logCPM) and the mean-variance relationship is modeled with precision weights (voom transform)

Limma Example

```
#From edgeR, Computes library size  
dge <- DGEList(counts=counts, group=group)  
#Gene Filtering  
counts<-counts[which(rowSums(cpm(counts)))>1),]  
dge <- calcNormFactors(dge) #TMM normalization
```

Limma Example

```
design<-model.matrix(~group)
#voom transform to calculate weights to
#eliminate mean-variance relationship
v<-voom(dge, design)
#use usual limma pipelines
fit<-lmFit(v,design)
fit<-eBayes(fit)
```

Limma Example

```
topTable(fit, coef=ncol(design))
```


##		logFC	AveExpr	t	P.Value	adj.P.Val	B
##	LINC01093	5.392080	0.8713944	7.178655	2.221795e-08	0.0002659532	9.050971
##	ZEB2	1.636165	6.4212073	7.042985	3.322131e-08	0.0002659532	8.750325
##	PGF	5.057301	-3.0391377	6.873776	5.499113e-08	0.0002934877	7.215693
##	LAYN	3.297688	-0.9927497	6.583385	1.312763e-07	0.0005254664	6.968214
##	FAM151B	3.510048	1.2897952	6.395061	2.315193e-07	0.0007413710	6.927036
##	C7orf61	2.431759	2.7351541	6.272193	3.355986e-07	0.0007874651	6.621596
##	TIPARP	2.073986	5.0819834	6.253512	3.551099e-07	0.0007874651	6.567031
##	DCUN1D3	2.253232	1.1217001	5.964492	8.528618e-07	0.0011379309	5.648099
##	CYP19A1	6.349196	-0.9436098	5.988963	7.917933e-07	0.0011379309	5.634653
##	PDGFC	3.546129	0.9321525	5.868816	1.140518e-06	0.0012173893	5.421180

Pathway analysis


After finding DEGs, look for correlated genes/networks and enriched pathway sets in the gene set using:

- ▶ Weighted gene coexpression network analysis (WGCNA)
- ▶ GSEA, GSVA, EnrichR, many more!!
- ▶ Qiagen Ingenuity Pathway Analysis (IPA)

Pathway analysis



Pathway analysis



TBSignatureProfiler Analysis

The TBSignatureProfiler was developed in the Johnson Lab in 2021 to profile new and existing TB gene expression signatures:

<https://bmcinfectdis.biomedcentral.com/articles/10.1186/s12879-020-05598-z>

TBSignatureProfiler Analysis


```
se_hivtb_2 <- se_hivtb[,  
  colData(se_hivtb)$disease_status != "TB-HIV-ART"]  
TBsigns <- TBsignatures[-12]  
ssgsea_res <- runTBsigProfiler(se_hivtb_2,  
  useAssay = "log_counts_cpm",  
  signatures = TBsigns,  
  algorithm = "ssGSEA",  
  combineSigAndAlgorithm = TRUE,  
  parallel.sz = 1)
```

Signature Heatmap:

```
# Colors for gradient
signatureHeatmap(ssgsea_res,
  name = "Heatmap of Signatures (ssGSEA)",
  signatureColNames = names(TBsigns),
  annotationColNames = c("disease_status"),
  scale = TRUE,
  split_heatmap = "none",
  showColumnNames = FALSE)
```


Signature Boxplots

```
signatureBoxplot(ssgsea_res, name="ssGSEA",
  signatureColNames = names(TBsigns)[c(62,77)],
  annotationColName = c("disease_status"))
```



Session info

```
sessionInfo()
```

```
## R version 4.5.2 (2025-10-31)
## Platform: aarch64-apple-darwin20
## Running under: macOS Tahoe 26.5.1
##
## Matrix products: default
## BLAS: /System/Library/Frameworks/Accelerate.framework/Versions/A/Frameworks/vecLib.framework/Versions/A/libBLAS.dylib
## LAPACK: /Library/Frameworks/R.framework/Versions/4.5-arm64/Resources/lib/libRlapack.dylib; LAPACK version 3.12.1
##
## locale:
## [1] en_US/en_US/en_US/C/en_US/en_US
##
## time zone: America/New_York
## tzcode source: internal
##
## attached base packages:
## [1] grid      stats4    stats      graphics  grDevices  utils      datasets
## [8] methods   base
##
## other attached packages:
## [1] umap_0.2.10.0          TBSignatureProfiler_1.20.0
## [3] ComplexHeatmap_2.24.1 DESeq2_1.48.2
## [5] edgeR_4.6.3            limma_3.64.3
## [7] SummarizedExperiment_1.38.1 Biobase_2.68.0
## [9] GenomicRanges_1.60.0    GenomeInfoDb_1.44.3
## [11] IRanges_2.42.0          S4Vectors_0.48.0
```